

Use Redis to Avoid Ticket Collision - Ticketing Platform

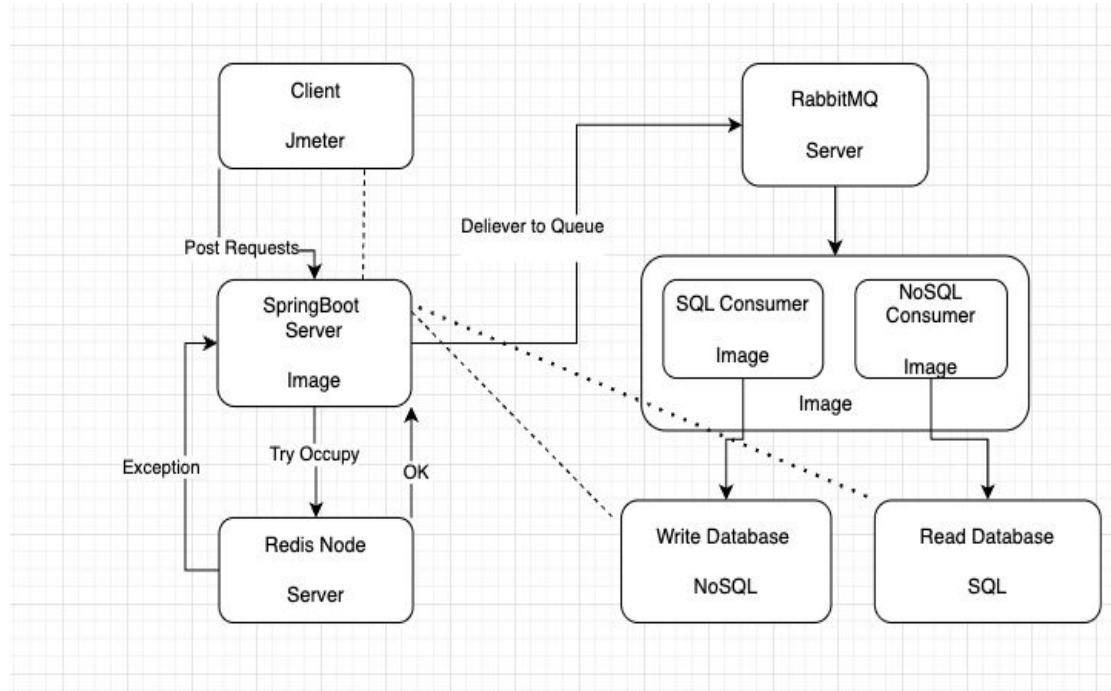
CS6650 Final Project – Zeyu Li

Project Design

- Simulate Ticket Purchasing Platform that users wait until start time to purchase the ticket directly
- Concurrent requests all together in the beginning phase, including duplication scenarios
- 100 Zones * A-Z Row * 30 seat
- Structure
 - **Venue:** venueld, **List of Zones**
 - **Event:** eventId, venueld, name, type, date
 - **Zone:** zoneld, **ticketPrice**, rowCount, colCount
 - **TicketInfo:** ticketId, venueld, eventId, zoneld, col, row, createdOn
 - **TicketCreation:** **status**

Project Structure

- Client: Jmeter
- Server: Java Spring Boot
 - Two Controllers
- Cache: Redis
- MQ: RabbitMQ
- NoSQL: Dynamodb
- SQL: MySQL



CQRS - Tickets/ and Query/

- Commands mutate state (POST /ticket → Redis → MQ → write store)
- Queries read state (GET /get/xxx → direct SQL and MySQL access)

Write Model (Commands)

- Redis + Lua for atomic seat reservation
- RabbitMQ to enqueue reservation events
- Designed DynamoDB for fast and idempotent ticket writes

Read Model (Queries)

- MySQL for real-time counts, revenue, and reporting
- Direct JDBC queries bypass the message queue

Redis - Zone → Row → Seat

- 3 Step Check to Provide Feedback if Duplicate Seat Purchase
- Bitmap: 1 bit per seat for $O(1)$ occupancy checks
- Zone Counter: tracks total available seats per zone
 - `event:{eventId}:zone:{zoneId}:remainingZoneSeats`
- Row Counter: tracks available seats per row
 - `event:{eventId}:zone:{zoneId}:row:{rowIndex}:remainingSeats`

Zone 1

A ₁	A ₂	A ₃	A ₃₀
B ₁	B ₂	B ₃	B ₃₀
C ₁	C ₂	C ₃	C ₃₀
<div>RowIndex × SeatsPerRow + ColIndex Pos = 3 × 30 + 3 = 93</div>					

Redis - Lua Scripts

- occupySeat.lua: GETBIT, DECR counters, SETBIT in one atomic script
- releaseSeat.lua: SETBIT=0 and INCR counters for rollback
- Single-threaded executions
 - Atomicity without external locks
 - Only one script at the same time
- Lettuce Connection Pool: max-active=8 to manage client concurrency

Deployment Attempt

EC2 deployment


- DynamoDB - IAM requirements
- RabbitMQ throughput problem

Local Docker Image Deployment

Overview


Totals

Queued messages (chart: last minute) (?)



Ready 10,581 msg
Unacknowledged 200 msg
Total 10,781 msg

Message rates (chart: last minute) (?)



Publish 0.00/s
Deliver 138/s
Redelivered 0.00/s
Acknowledge 138/s

Global counts (?)

Connections: 2

Channels: 40

Exchanges: 9

Queues: 3

Consumers: 15

Nodes

Name	File descriptors (?)	Socket descriptors (?)	Erlang processes	Memory	Disk space	Uptime	Type
rabbit@ip-172-31-20-96	25 65535 available	3 58889 available	352 1048576 available	74MB 381MB high watermark	3.1GB 48MB low watermark	3h 5m	Disc Stats ...

Ports and contexts

Local Test Thread Setting

MySQL (HikariCP) Connection Pool

- Minimum Idle Connections: 400
- Maximum Pool Size: 2000

Redis (Lettuce) Connection Pool

- Max Active Connections: 8
- Max Idle Connections: 8
- Min Idle Connections: 0
- Max Wait Time: infinite

RabbitMQ Listener Container

- Initial Concurrent Consumers: 100
- Maximum Concurrent Consumers: 300
- Prefetch Count: 300 messages per consumer

DynamoDB HTTP Client Pool

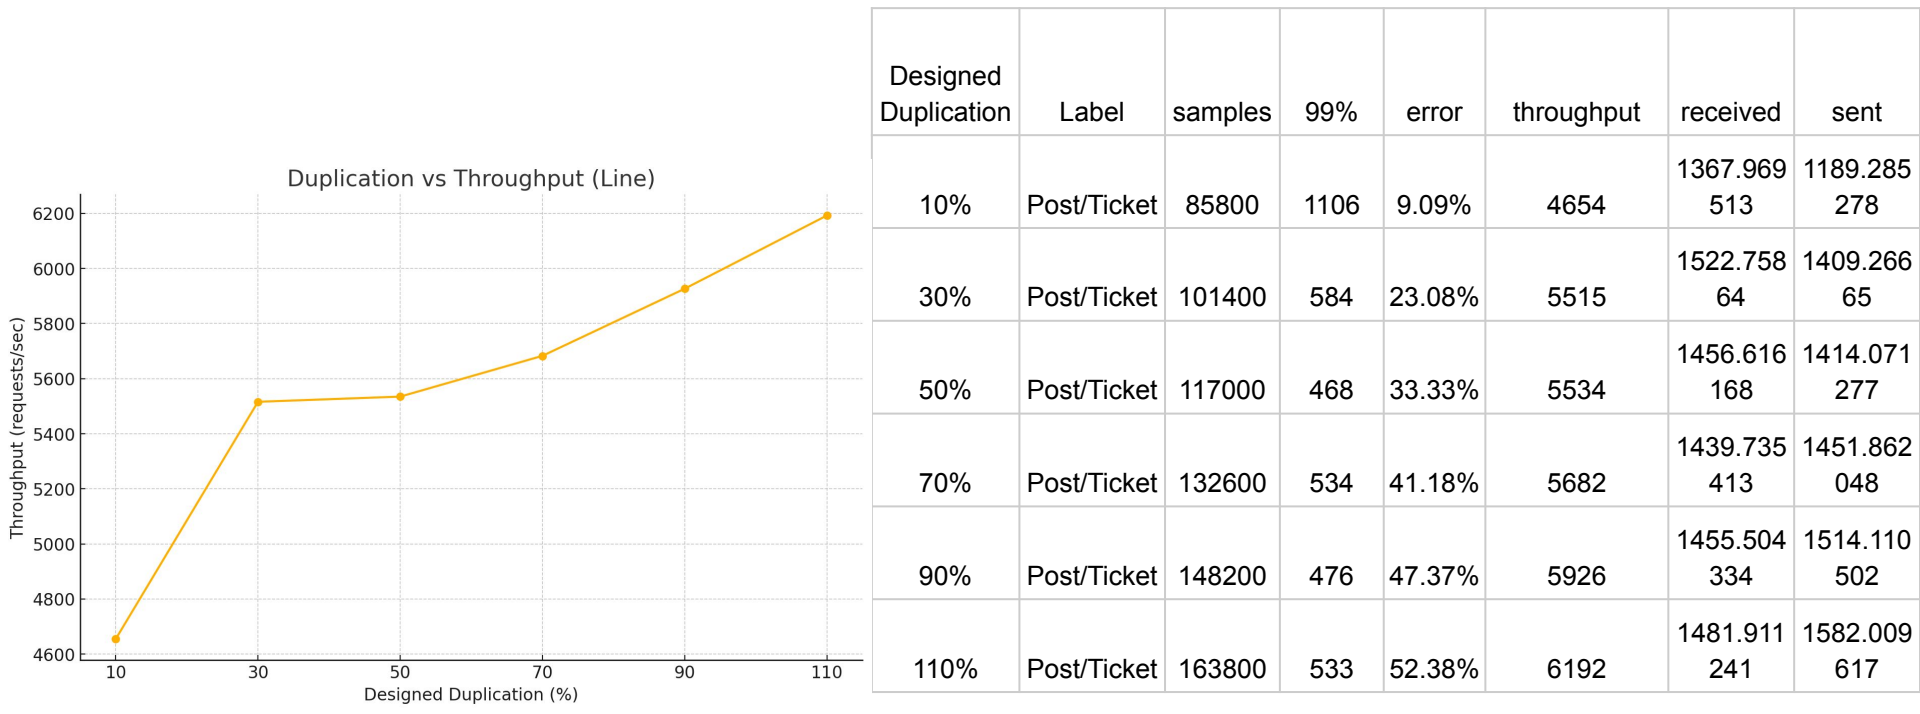
- Max Connections: 3000
- Connection Timeout: 10 seconds
- Socket (Read/Write) Timeout: 30 seconds

JMeter Load Test

- Client Threads: 1000 concurrent threads

Test Result - Increase duplication

At the same setting of threads and environment, increase the duplication rate in the dataset to mimic more buyers for the same ticket.



Future Work

- Redis Sharding:
 - add a second node and hash by zoneid to run two Lua scripts in parallel
 - higher reservation throughput
- Read-Only Cache:
 - cache new ticket details in a separate read-only Redis instance for instant “My Tickets” lookups without impacting write load
- Transparent Routing & Holds:
 - Support dynamic hold durations
 - Support a transparent tiered processing speed model

Thank you!